

(2) Modeling

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as scs
import seaborn as sns
import plotly.express as px
import statsmodels.api as sm
```

```
In [2]: df = pd.read_csv('bam_data.csv')
df.head()
```

Out[2]:

	approach_vertical	vertical_jump	three_quarter_court_sprint	four_way_agility	reaction_shuttle	t
0	33.5	28.5	3.376	11.471	3.669	
1	30.5	21.5	3.486	12.114	3.355	
2	37.0	31.0	3.230	12.036	3.562	
3	29.0	23.0	3.370	12.509	3.173	
4	31.0	26.0	3.389	12.724	3.316	

```
In [3]: df.describe()
```

Out[3]:

	approach_vertical	vertical_jump	three_quarter_court_sprint	four_way_agility	reaction_shuttle	t
count	1059.000000	1059.000000	1059.000000	1059.000000	1059.000000	1059.000000
mean	31.829615	25.860157	3.467047	12.247189	3.505240	3.505240
std	3.423395	3.065653	0.335502	0.652726	0.273780	0.273780
min	19.000000	14.000000	2.950000	10.359000	2.914000	2.914000
25%	30.000000	24.000000	3.339500	11.806000	3.348000	3.348000
50%	31.829615	25.860157	3.424000	12.244000	3.492000	3.492000
75%	34.000000	28.000000	3.537500	12.658000	3.634000	3.634000
max	43.500000	38.000000	9.954000	14.775000	6.759000	6.759000

6) - Find best model

```
In [4]: # https://scikit-learn.org/stable/modules/tree.html#classification
# https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_s
caling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py
# https://towardsdatascience.com/data-science-mistakes-to-avoid-data-lea
kage-e447f88aae1c
### - FIXED DATA LEAKAGE
```

1) Split train/test data

```
In [5]: from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler, normalize
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn import tree
from sklearn.datasets import load_digits
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.ensemble import RandomForestRegressor
from sklearn.datasets import make_regression
from sklearn.metrics import r2_score

y = df['bamscore']
X = df.drop('bamscore',axis=1) #drop bamscorerank too once I added so no
data leakage
# y is bamscorerank, x is everything else
# now xtrain/split will take those two variables and create two datasets
from it

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=0, stratif
y = None)
```

```
In [6]: train_test_split(y, shuffle=False)
```

```
Out[6]: [0      2003.0
         1      1865.0
         2      2005.0
         3      1902.0
         4      1903.0
         ...
        789     1916.0
        790     1852.0
        791     1593.0
        792     1814.0
        793     2114.0
        Name: bamscore, Length: 794, dtype: float64, 794      1843.000000
        795     1968.000000
        796     1965.000000
        797     1950.000000
        798     1902.000000
        ...
       1054     1917.000000
       1055     2029.000000
       1056     1890.976326
       1057     1890.976326
       1058     1890.976326
        Name: bamscore, Length: 265, dtype: float64]
```

2) Normalize + Scale Data to adjust for outliers

- going to use min/max normalization to make all points between 0-1
- meaning, make data all within 0-1 range. 5.5 on 0-10 scale would be .55 on normalized scale

Why are we normalizing data?

to make data cleaner, sometimes gets messed up on backend and removes unstructured data

```
In [7]: y_train=(y_train-y_train.min())/(y_train.max()-y_train.min())
        y_test=(y_test-y_test.min())/(y_test.max()-y_test.min())
```

```
In [8]: normalize(X_train, copy = False)
        normalize(X_test, copy = False)
        df_train = pd.DataFrame(data = np.concatenate((X_train.values, y_train.v
        alues.reshape(-1,1)),axis = 1)
                                , columns = list(X_train.columns) + ["bamscore"
        ])
```

```
In [9]: df_train
```

```
Out[9]:
```

	approach_vertical	vertical_jump	three_quarter_court_sprint	four_way_agility	reaction_shuttle
0	0.127644	0.110904	0.015016	0.049245	0.015585
1	0.139731	0.113526	0.015712	0.052135	0.014799
2	0.173806	0.130354	0.016736	0.060029	0.016363
3	0.108274	0.098607	0.013093	0.046635	0.012846
4	0.159682	0.126732	0.018128	0.065885	0.017337
...
736	0.149692	0.128307	0.015715	0.054516	0.016851
737	0.181453	0.144703	0.013625	0.056108	0.015288
738	0.127032	0.098270	0.016855	0.060357	0.015843
739	0.143284	0.115483	0.013507	0.050077	0.013610
740	0.155346	0.168662	0.015299	0.051544	0.015202

741 rows × 13 columns

```
In [10]: # https://scikit-learn.org/stable/auto\_examples/preprocessing/plot\_all\_scaling.html#sphx-glr-auto-examples-preprocessing-plot-all-scaling-py
```

3) Feature Importance + Feature Selection

```
In [11]: from sklearn.datasets import load_digits
from sklearn.feature_selection import SelectKBest, f_regression
```

```
In [12]: X_train
```

```
Out[12]:
```

	approach_vertical	vertical_jump	three_quarter_court_sprint	four_way_agility	reaction_shuttle
899	0.127644	0.110904	0.015016	0.049245	0.015585
635	0.139731	0.113526	0.015712	0.052135	0.014795
310	0.173806	0.130354	0.016736	0.060029	0.016365
961	0.108274	0.098607	0.013093	0.046635	0.012845
723	0.159682	0.126732	0.018128	0.065885	0.017337
...
1033	0.149692	0.128307	0.015715	0.054516	0.016851
763	0.181453	0.144703	0.013625	0.056108	0.015285
835	0.127032	0.098270	0.016855	0.060357	0.015845
559	0.143284	0.115483	0.013507	0.050077	0.013610
684	0.155346	0.168662	0.015299	0.051544	0.015205

741 rows × 12 columns

```
In [13]: # want to fit and transform the model on training x and y data
```

```
In [14]: #skb = SelectKBest(f_regression, k=8)
skb = SelectKBest(score_func=f_regression, k=5)
fit = skb.fit(X_train, y_train)
features = fit.transform(X_train)
#X_new = SelectKBest(chi2, k=20).fit_transform(X, y)
```

```
In [15]: mask = fit.get_support()
X_train[X_train.columns[mask]]
```

Out[15]:

	approach_vertical	vertical_jump	wingspan	weight	hand_width
899	0.127644	0.110904	0.321202	0.771723	0.037665
635	0.139731	0.113526	0.338029	0.726981	0.038412
310	0.173806	0.130354	0.370615	0.675797	0.039617
961	0.108274	0.098607	0.311287	0.781892	0.034802
723	0.159682	0.126732	0.372591	0.679282	0.045623
...
1033	0.149692	0.128307	0.332649	0.751786	0.039205
763	0.181453	0.144703	0.328452	0.739132	0.039047
835	0.127032	0.098270	0.342747	0.733430	0.043143
559	0.143284	0.115483	0.322924	0.761332	0.041702
684	0.155346	0.168662	0.343981	0.723470	0.037727

741 rows × 5 columns

```
In [16]: print(mask)
print(features)
print(features.shape)
```

```
[ True  True False False False  True False False  True False False  True
]
[[0.12764397 0.11090378 0.32120246 0.77172292 0.03766544]
 [0.13973142 0.11352561 0.33802858 0.72698095 0.03841234]
 [0.17380558 0.13035419 0.37061485 0.675797 0.03961745]
 ...
 [0.12703203 0.09827006 0.34274681 0.73343023 0.04314295]
 [0.14328437 0.11548293 0.32292448 0.76133189 0.04170217]
 [0.15534627 0.16866167 0.34398103 0.72346978 0.03772695]]
(741, 5)
```

```
In [17]: ### OLS Regression (Ordinary Least Squares)
### Use this because it's the most simple baseline model
### Model gives best approximate of true population regression line.
### The principle of OLS is to minimize the square of errors
```

```
In [18]: import statsmodels.formula.api as smf

formula = "bamscore ~ approach_vertical + vertical_jump + four_way_agility"
lm = smf.ols(formula = formula, data = df_train).fit()
print(lm.summary())
print("After selecting best 3 features:", features.shape)
```

OLS Regression Results

```

=====
Dep. Variable:          bamscore    R-squared:
0.619
Model:                  OLS         Adj. R-squared:
0.617
Method:                 Least Squares    F-statistic:
398.5
Date:                   Mon, 25 Apr 2022    Prob (F-statistic):      8.
99e-154
Time:                   12:02:45    Log-Likelihood:
714.13
No. Observations:      741    AIC:
-1420.
Df Residuals:          737    BIC:
-1402.
Df Model:               3
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.0
25	0.975]				
Intercept	0.5078	0.034	15.015	0.000	0.4
41 approach_vertical	3.9594	0.365	10.837	0.000	3.2
42 vertical_jump	3.3054	0.414	7.976	0.000	2.4
92 four_way_agility	-16.3254	0.752	-21.718	0.000	-17.8
01	-14.850				

```

=====
Omnibus:                72.121    Durbin-Watson:
2.015
Prob(Omnibus):          0.000    Jarque-Bera (JB):
219.234
Skew:                   -0.459    Prob(JB):
2.48e-48
Kurtosis:               5.502    Cond. No.
229.
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

After selecting best 3 features: (741, 5)


```

In [19]: ##### df = samplesize - # of variables +1 --> number of independent observations
##### constant is intercept in regression line and tells us avg value of omitted variables and noise in model
##### coeff term = slope aka rise over run, if x increases by 1 , and coefficient is .7, y increased by .7
##### standard error = std
##### Standard error shows the sampling variability of these parameters.  $\sigma^2$  is equal to residual sum squares
##### remember  $\sigma^2$  in first param and its the numerator in second param

##### H0 : B2 = 0 ( variable X has no influence on Y)
##### Ha : B2  $\neq$  0 (X has significant impact on Y)
##### b1 ~ N(B1,  $\sigma_{b1}^2$ ) B1 is true mean of b1
##### - b1 is param
##### b2 ~ N(B2 ,  $\sigma_{b2}^2$ ) B2 is true mean of b2
##### - b2 is param

##### p value - reject null <0.05 or fail to reject if >0.05. if 0 t is probably high

##### R2 is the coefficient of determination that tells us that
##### how much percentage variation independent variable can be explained by independent variable.

# F stat - prob f stat > f stat given = 0 means reject null

##### https://www.geeksforgeeks.org/interpreting-the-results-of-linear-regression-using-ols-summary/

```

Add features to try and improve r^2 + find important features

```
In [20]: import statsmodels.formula.api as smf
formula = "bamscore ~ approach_vertical + vertical_jump + reach + weight
+ body_comp + four_way_agility"
lm = smf.ols(formula = formula, data = df_train).fit()
print(lm.summary())
print(mask)
print(features)
print(features.shape)
#https://www.datatechnotes.com/2021/02/selectbest-feature-selection-example-in-python.html
```

OLS Regression Results

```

=====
Dep. Variable:          bamscore    R-squared:
0.623
Model:                  OLS        Adj. R-squared:
0.620
Method:                 Least Squares    F-statistic:
202.5
Date:                   Mon, 25 Apr 2022    Prob (F-statistic):      5.
79e-152
Time:                   12:02:45    Log-Likelihood:
718.80
No. Observations:      741    AIC:
-1424.
Df Residuals:          734    BIC:
-1391.
Df Model:               6
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.0
25	0.975]				
Intercept	0.9057	0.167	5.412	0.000	0.5
77 approach_vertical	3.8064	0.370	10.301	0.000	3.0
81 vertical_jump	3.1563	0.417	7.568	0.000	2.3
38 reach	-0.0871	0.168	-0.519	0.604	-0.4
16 weight	-0.3362	0.125	-2.699	0.007	-0.5
81 body_comp	-0.0270	0.158	-0.171	0.864	-0.3
37 four_way_agility	-17.6547	0.904	-19.536	0.000	-19.4
29	-15.881				

```

=====
Omnibus:                73.223    Durbin-Watson:
2.011
Prob(Omnibus):           0.000    Jarque-Bera (JB):
227.170
Skew:                    -0.460    Prob(JB):
4.68e-50
Kurtosis:                5.552    Cond. No.
360.
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[ True  True False False False  True False False  True False False  Tru
e]

```

```
[[0.12764397 0.11090378 0.32120246 0.77172292 0.03766544]
 [0.13973142 0.11352561 0.33802858 0.72698095 0.03841234]
 [0.17380558 0.13035419 0.37061485 0.675797 0.03961745]
 ...
 [0.12703203 0.09827006 0.34274681 0.73343023 0.04314295]
 [0.14328437 0.11548293 0.32292448 0.76133189 0.04170217]
 [0.15534627 0.16866167 0.34398103 0.72346978 0.03772695]]
(741, 5)
```

```
In [21]: # body comp and weight >0.05 - Ideally just drop them
```

Iteration - I wanted to take out because it is hard to measure already, now we know we can take it out because it's really throwing off the data and martin said it does not matter much, but model does not know that

```
In [22]: import statsmodels.formula.api as smf
formula = "bamscore ~ approach_vertical + vertical_jump + reach + four_w
ay_agility"
lm = smf.ols(formula = formula, data = df_train).fit()
print(lm.summary())
sns.lmplot(x="four_way_agility", y="bamscore" ,data=df_train)
```


```

                                OLS Regression Results
=====
Dep. Variable:                  bamscore    R-squared:
0.620
Model:                          OLS        Adj. R-squared:
0.618
Method:                        Least Squares    F-statistic:
299.8
Date:                          Mon, 25 Apr 2022    Prob (F-statistic):
16e-153
Time:                          12:02:45    Log-Likelihood:
715.14
No. Observations:              741    AIC:
-1420.
Df Residuals:                  736    BIC:
-1397.
Df Model:                      4
Covariance Type:              nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.0
25      0.975]
-----
-----
Intercept                    0.4708      0.043      11.028      0.000      0.3
87      0.555
approach_vertical            3.9304      0.366      10.749      0.000      3.2
13      4.648
vertical_jump                3.2492      0.416       7.810      0.000      2.4
32      4.066
reach                       0.1890      0.133       1.418      0.157     -0.0
73      0.451
four_way_agility           -16.9350      0.866     -19.566      0.000     -18.6
34     -15.236
=====
=====
Omnibus:                      69.439    Durbin-Watson:
2.018
Prob(Omnibus):                0.000    Jarque-Bera (JB):
228.944
Skew:                         -0.410    Prob(JB):
1.93e-50
Kurtosis:                     5.597    Cond. No.
283.
=====
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.

```

Out[22]: <seaborn.axisgrid.FacetGrid at 0x7ff491df7fd0>



Iteration - features I think are most important from pairplot in EDA

```
In [23]: import statsmodels.formula.api as smf
formula = "bamscore ~ three_quarter_court_sprint + four_way_agility + re
action_shuttle + vertical_jump + approach_vertical"
lm = smf.ols(formula = formula, data = df_train).fit()
print(lm.summary())
```


OLS Regression Results

```

=====
=====
Dep. Variable:          bamscore    R-squared:
0.697
Model:                  OLS         Adj. R-squared:
0.695
Method:                 Least Squares    F-statistic:
337.9
Date:                  Mon, 25 Apr 2022    Prob (F-statistic):          1.
01e-187
Time:                  12:02:46    Log-Likelihood:
799.16
No. Observations:      741    AIC:
-1586.
Df Residuals:          735    BIC:
-1559.
Df Model:              5
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				

Intercept	0.5582	0.030	18.313	0.000
0.498 0.618				
three_quarter_court_sprint	-7.9533	1.923	-4.136	0.000
-11.728 -4.179				
four_way_agility	-6.0437	1.019	-5.932	0.000
-8.044 -4.044				
reaction_shuttle	-35.3052	2.773	-12.732	0.000
-40.749 -29.861				
vertical_jump	3.2517	0.370	8.788	0.000
2.525 3.978				
approach_vertical	4.4234	0.328	13.489	0.000
3.780 5.067				

```

=====
=====
Omnibus:              76.718    Durbin-Watson:
2.017
Prob(Omnibus):        0.000    Jarque-Bera (JB):
457.485
Skew:                 -0.204    Prob(JB):              4.
55e-100
Kurtosis:             6.828    Cond. No.
955.
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [24]: import statsmodels.formula.api as sm
model = sm.glsar(formula = formula, data = df_train)

lm = model.fit()
print(lm.summary())
```

```

GLSAR Regression Results
=====
Dep. Variable:          bamscore    R-squared:
0.697
Model:                  GLSAR      Adj. R-squared:
0.695
Method:                 Least Squares    F-statistic:
338.1
Date:                   Mon, 25 Apr 2022    Prob (F-statistic):          1.
10e-187
Time:                   12:02:46    Log-Likelihood:
798.14
No. Observations:      740    AIC:
-1584.
Df Residuals:          734    BIC:
-1557.
Df Model:               5
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				
Intercept	0.5579	0.030	18.301	0.000
three_quarter_court_sprint	-7.9910	1.923	-4.156	0.000
four_way_agility	-5.9933	1.020	-5.877	0.000
reaction_shuttle	-35.4437	2.776	-12.769	0.000
vertical_jump	3.2400	0.370	8.753	0.000
approach_vertical	4.4347	0.328	13.518	0.000

```

=====
=====
Omnibus:               76.581    Durbin-Watson:
2.013
Prob(Omnibus):         0.000    Jarque-Bera (JB):
458.040
Skew:                  -0.201    Prob(JB):              3.
45e-100
Kurtosis:              6.833    Cond. No.
955.
=====
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.

```

```
In [25]: ypred = lm.predict(X_test)
```

```
In [26]: r2_score(y_test, ypred)
```

```
Out[26]: 0.3491785846088724
```

```
In [27]: # DO LM PLOT to see trend in data with respect to bamscore
```

```
In [28]: # Random Forest  
## Going to try random forrest because we had over fitting in models above  
## Random forrest corrects for high variance/overfitting by using many decision trees  
  
#Steps:  
# -create bootstrapped dataset with subset of variables  
# -fit decision tree  
# -repeat and tally predictions
```

```
In [29]: from sklearn.ensemble import RandomForestRegressor  
  
# create regressor object  
regressor = RandomForestRegressor(n_estimators = 100, random_state = 0)  
cols = ["four_way_agility", "reaction_shuttle", "three_quarter_court_sprint", "vertical_jump", "approach_vertical"]  
# fit the regressor with x and y data  
rf = regressor.fit(X_train, y_train)  
rf.score(X_test, y_test)  
# score if fit and test on full data frame
```

```
Out[29]: 0.5295593565201384
```

```
In [30]: #x_train[cols]  
#x_test[cols]  
  
regressor = RandomForestRegressor(n_estimators = 100, random_state = 0)  
cols = ["four_way_agility", "reaction_shuttle", "three_quarter_court_sprint", "vertical_jump", "approach_vertical"]  
# fit the regressor with x and y data  
rf = regressor.fit(X_train[cols], y_train)  
rf.score(X_test[cols], y_test)  
# score if fit and test on just [cols]
```

```
Out[30]: 0.5382597305545691
```

```
In [31]: # If Condition Number is too high I'll run into problems in matrix because program may not catch small errors
```

```
In [32]: # 5) Iterative modeling process  
### What models are appropriate  
### Compare Models  
### Find which performance metrics to use and adjust to make the model better.
```

```
In [33]: #### Assumptions to test in model:  
#### - Combine Tests more important than physical measurments. Weigh com  
bine tests as double.  
#### - just tests, just measurments, 1:1 test/measurements, hypothesis -  
2:1 test/measurment
```

```
In [34]: #r2 = model isn't very accurate predicting bamscore with .453 bamscore
```

```
In [35]: # Least Sum Squares to figure out best fit line and figure out most corr  
related features
```